

# Tidying Data in R

Bas Machielsen

Utrecht University

December 29, 2019

# What is tidy data?

- In order for your data to be ready for analysis, it needs to be **tidy**.
- Tidy data is defined as:  
*data sets that are arranged such that each variable is a column and each observation (or case) is a row.*
- I will attempt to explain how to put this principle in practice using a set of packages called the **tidyverse**, but strictly speaking, only the `tidyr` package is required.
- You can download the `.Rmd` and `.csv` files from my **github page**, and follow along!

# What is (un)tidy data?

*All tidy data are similar to each other. All untidy data are untidy in their own way.*

- Although tidy data shares the properties mentioned in the previous slide, untidy data refers (broadly speaking) to unstructured data.
- An example (from [www.tidyverse.org](http://www.tidyverse.org)):

Table 1: This is what an untidy dataset looks like

name	treatmenta	treatmentb
John Smith	NA	18
Jane Doe	4	1
Mary Johnson	6	7

# An example of untidy data

- Is each observation in a row?
- Is each variable in a column?
- What are the variables here?

Table 2: This is also what an untidy dataset looks like

treatment	John.Smith	Jane.Doe	Mary.Johnson
a	NA	4	6
b	18	1	7

## Example: from untidy to tidy data

- The variables are (i) the treatment, (ii) the subject, and (iii) the 'score' corresponding to each subject-treatment observations.

```
preg %>%  
  pivot_longer(treatmenta:treatmentb,  
               names_to = "treatment", values_to = "score") %>%  
  mutate(treatment = gsub("treatment", "", treatment)) %>%  
  arrange(name, treatment) %>%  
  kable(caption = "Test", booktabs = TRUE,  
        row.names = FALSE) %>%  
  kable_styling(latex_options = "striped")
```

## Example: from untidy to tidy data

- The former chunk of code transforms the data in the first table to the following table:

Table 3: This is tidy data

name	treatment	score
Jane Doe	a	4
Jane Doe	b	1
John Smith	a	NA
John Smith	b	18
Mary Johnson	a	6
Mary Johnson	b	7

# Transforming the data

- The key command in the former chunk of code is `pivot_longer`. With the arguments `data`, `columns`, `names_to` and `values_to`.
- All the other functions are essentially layout changes, not fundamental transformations.
- The command `pivot_longer` basically transforms the dataset from:

Table 4: This is what an untidy dataset looks like

name	treatmenta	treatmentb
John Smith	NA	18
Jane Doe	4	1
Mary Johnson	6	7

# Transforming the data

- to :

```
pivot_longer(data = preg, cols = 2:3, names_to = "treatment",
             values_to = "score") %>%
  kable(caption = "This is what a tidy dataset looks like",
        booktabs = TRUE, row.names = FALSE) %>%
  kable_styling(latex_options = "striped")
```

Table 5: This is what a tidy dataset looks like

name	treatment	score
John Smith	treatmenta	NA
John Smith	treatmentb	18
Jane Doe	treatmenta	4
Jane Doe	treatmentb	1
Mary Johnson	treatmenta	6
Mary Johnson	treatmentb	7



- The best way to see how tidy data works is by using examples. In this lecture, I will demonstrate extensively how to create tidy data that is ready for analysis using two examples.
- First, I will show how to combine and tidy data on economic development from the World Bank.
- Afterwards, I show how to combine and tidy data from Amadeus, a commercial database.
- On the fly, I will demonstrate how simple it is to analyze or create graphs with tidy data.
- Finally, I will use other examples to show how to tidy more messy data, inspired by examples from the **tidyverse website**

# World Bank Dataset

- Let us now download some data from the World Bank database. I extract the 55 basic World Development Indicators from **here** for 20 countries.
- Importing the data..

```
worldbank <- read_csv("wb.csv")
```

```
dim(worldbank)
```

```
## [1] 1105 14
```

- The data has 1105 observations and 14 variables. That is not at all what we wanted. What do the data look like?

# World Bank Dataset

```
kable(worldbank[1:4,c(1,2,4,5)],  
      caption = "Work Bank Data - Untidy!",  
      booktabs = TRUE, row.names = FALSE) %>%  
kable_styling(latex_options = "striped")
```

Table 6: Work Bank Data - Untidy!

Country Name	Country Code	Series Code	2009 [YR2009]
Argentina	ARG	SP.ADO.TFRT	63.1896
Argentina	ARG	NV.AGR.TOTL.ZS	5.27362346890139
Argentina	ARG	ER.H2O.FWTL.ZS	..
Argentina	ARG	SH.STA.BRTC.ZS	97.9

- Very untidy dataset! NA observations are entered as . . , and variable names require an extensive definition before you know what they mean.
- Furthermore, years are not notated straightforwardly, and the data violates the tidy data principles.

# What to do?

- First, let us try to save the variable names and series code to another dataset (step 1), and delete the names from the worldbank dataset (step 2).

## *#Step 1*

```
wbnames <- cbind(unique(worldbank[,3]),unique(worldbank[,4])) %>%  
  na.omit()
```

## *#Step 2*

```
worldbank <- worldbank[1:1100,-3]
```

# Almost there..

- Now, let's try to unpivot the columns containing the years..

```
worldbank <- pivot_longer(worldbank, 4:13,  
  names_to = "years",  
  values_to = "value")
```

- This is what the data looks like right now.

Table 7: The data then looks like this.

Country Code	Series Code	years	value
ARG	SP.ADO.TFRT	2009 [YR2009]	63.1896
ARG	SP.ADO.TFRT	2010 [YR2010]	63.3154
ARG	SP.ADO.TFRT	2011 [YR2011]	63.4412
ARG	SP.ADO.TFRT	2012 [YR2012]	63.567

# Final steps

- Let's now clean the “years” variable (step 1), set the .. observations to NA (step 2), and “widen” the data so that every separate indicator gets a column (step 3).
- Step 3 is done using the `pivot_wider` command: it widens the data and shortens the number of observations, because it transfers information from rows to columns.

## *#Step 1*

```
worldbank <- worldbank %>%  
  mutate(years = substring(years, 1, 4))
```

## *#Step 2*

```
worldbank$value[worldbank$value == ".."] <- NA
```

## *#Step 3*

```
worldbank <- pivot_wider(data = worldbank,  
  names_from = `Series Code`,  
  values_from = `value`)
```

- `pivot_wider`, takes three arguments: the data, from which column it should take the new column names (`names_from`), and from which column it should take the values belonging to the corresponding cels (`values_from`).
- This is what it looks like!

Table 8: Tidy Data

Country Name	Country Code	years	SP.ADO.TFRT	NV.AGR.TOTL.ZS
Argentina	ARG	2009	63.1896	5.27362346890139
Argentina	ARG	2010	63.3154	7.13216745078964
Argentina	ARG	2011	63.4412	6.99873377022519
Argentina	ARG	2012	63.567	5.7817442068501

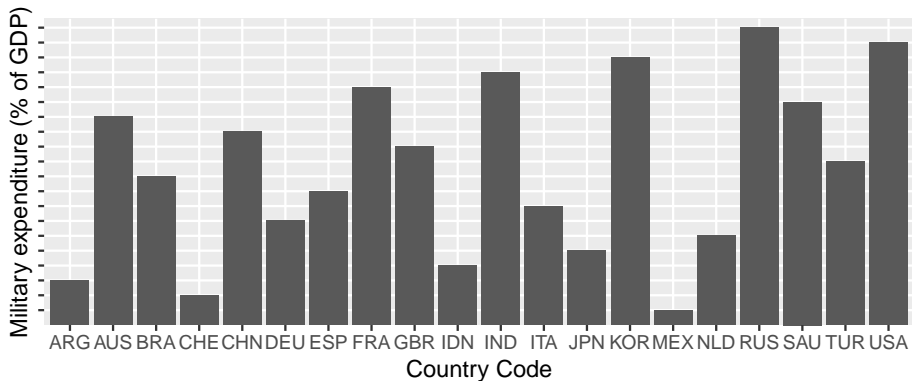
# Military Spending

- Suppose I want to know the military expenditures as a percentage of GDP MS.MIL.XPND.GD.ZS for each country in the dataset in 2015.
- I can use the `wbnames` information we've saved to add labels to the data.

```
worldbank %>%  
  filter(years == "2015") %>%  
  ggplot(aes(x = `Country Code`, y = `MS.MIL.XPND.GD.ZS`)) +  
  geom_col() +  
  labs(y = as.character(wbnames[match("MS.MIL.XPND.GD.ZS",  
                                     wbnames[,2]),1])) +  
  theme(axis.text.y = element_blank())
```



# Military Spending




# Fertility Rate

- Suppose now you want to track the evolution of the fertility rate over the last 10 years of some countries in the dataset.

```
worldbank %>%  
  filter(years < 2018,  
         `Country Code` == "CHE" |  
         `Country Code` == "GBR"  
         | `Country Code` == "RUS" |  
         `Country Code` == "SAU") %>%  
  ggplot(aes(x = years, y = `SP.DYN.TFRT.IN`,  
            group = `Country Code`,  
            color = `Country Code`)) +  
  geom_line() +  
  labs(y = as.character(wbnames[match("SP.DYN.TFRT.IN",  
                                       wbnames[,2]),1])) +  
  theme(axis.text.y = element_blank(),  
        axis.ticks.y = element_blank())
```

# Fertility Rate



## Another untidy dataset..

- Now, I will download some data from Amadeus, a commercial database collected information about firms across the world.
- I collect data from 2000 European companies on their financial information, the number of employees, and the number of directors/managers.

```
amadeus <- read.csv("Amadeus_Export_1.csv", fileEncoding = "UCS-2LE")
```

```
dim(amadeus)
```

```
## [1] 2000 68
```

# What does the data look like?

- These are some of the variables..

Table 9: Variable names

---

Mark
Company.name
City
Country.ISO.code
NACE.code
Cons..code
Last.year
Operating.revenue..Turnover..th.EUR.2019
Operating.revenue..Turnover..th.EUR.2018

---

# Tidying the Amadeus data

- The data look very untidy. Column names contain variables and years, and there are various superfluous columns, such as `Mark`, and last year. NA observations are coded as `n.a.` but are not recognized by `are` as such.
- The first thing we need to do is to change the NA obs to functioning NA obs, and remove some superfluous columns

## *#Step 1*

```
amadeus[amadeus == "n.a."] <- NA
amadeus[amadeus == "n.s."] <- NA
amadeus <- amadeus[, -c(1, 5:7)]
```

# Using pivot to transform the data

- Next, we want to convert the data into a **long** format, using `pivot_longer`.

```
amadeus <- amadeus %>%  
  mutate_all(as.character)  
  
amadeus <- pivot_longer(data = amadeus,  
                        cols = 4:64,  
                        names_to = "variable",  
                        values_to = "value")
```

# Using pivot to transform the data

- This is what the data looks like now:

Table 10: The amadeus dataset

x
Company.name
City
Country.ISO.code
variable
value



# String correction

- Now, we correct the strings using the following steps:

```
#First, remove "th.EUR" from the string
amadeus$variable <- sub("\\.th.EUR.", "",
                        amadeus$variable)

#Extract the year and put it into a new variable
amadeus$year <- as.numeric(
  str_extract(amadeus$variable, "[0-9]+"))

#Include only letters from the alphabet as variable names
amadeus$variable <- sapply(
  str_extract_all(
    amadeus$variable, "[aA-zZ]+"),
  paste, collapse = "_")
```

# Pivoting...

- Finally, we can expand the table to adhere to the principal of one obs. per row, one variable per column.

```
amadeus <- pivot_wider(data = amadeus, names_from = variable,  
                      values_from = value)
```

*#There is only one problem, No. of directors/managers  
#violates the tidy data principles.*

```
nodir <- amadeus %>%  
  filter(is.na(year))  
amadeus <- amadeus %>%  
  filter(!is.na(year))
```

```
amadeus <- merge(amadeus[,1:16],  
                nodir[,c(1,17)],  
                by = 1)
```

```
amadeus[,5:17] <- data.frame(sapply(amadeus[,5:17],  
                                   function(x) as.numeric(x)))
```

- Now, the data looks like this:

Table 11: Tidy data

Company.name	year	Operating_revenue_Turnover
@JCG-GMAO-CONSULTING	2019	88
@JCG-GMAO-CONSULTING	2018	133
@JCG-GMAO-CONSULTING	2017	62
@JCG-GMAO-CONSULTING	2016	51
@JCG-GMAO-CONSULTING	2015	NA

- All other variables are located in the remaining columns.

# Some analyses

- It seems that cleaning this type of data already required a fair amount of work, including some programming that requires knowledge of regular expressions, and some fairly non-standard transformation.
- In the remainder, I show some other examples of `pivot_longer` and `pivot_wider` in non-standard situations, so you do not get stuck when coming across these types of data.